

AI-SPM — Feature-Wise Sprint

Four-sprint delivery plan · Feature-wise implementation backlog

01

Sprint 1 — Company Onboarding & First Connected Device

1 Aug - 14 Aug · Labels: SAAS, ONBOARDING, SUPER ADMIN, ENDPOINT AGENT

#	JIRA TITLE	DESCRIPTION
1	Security service foundation (company data model, health checks)	Set up the central service that stores companies, users, devices and activity, with health checks so the team can confirm it is running.
2	Create company workspace, first admin user and install token	When a company joins, automatically create its private workspace, its first admin account and the install token used to connect devices.
3	Company admin login and secure session	A company admin can sign in with email and password, stay signed in securely, and sign out.
4	Company self-registration page (company name, admin email, password)	A new customer registers without contacting sales by entering their company name, admin email and password.
5	Email verification and first-login welcome onboarding	The admin confirms their email address, then sees a short welcome guide explaining the first steps to protect their company.
6	Public entry point with safe routing and request limits	Put all customer traffic behind a single protected entry point that limits abusive request volumes.
7	Super Admin: list, suspend and reactivate customer companies	The Super Admin sees every customer company and can suspend or reactivate an account, for example for non-payment.
8	Company's first AI protection policy (allow / mask / block)	The company admin chooses what happens when sensitive information is found in an AI prompt: allow it, mask it or block it.
9	Agent application skeleton that runs on a Linux machine	Build the endpoint application that runs on an employee computer, reads its configuration and stays running in the background.

#	JIRA TITLE	DESCRIPTION
10	Device enrolment service: a device joins a company using its install token	The service accepts a device's install token, confirms which company it belongs to and issues its identity.
11	Agent registers itself with the company install token	On first start the agent enrolls itself with the company install token so it appears in that company's device list.
12	Agent reports it is alive so the company sees Online / Offline	The agent checks in regularly, and a device that stops checking in is shown as Offline in the console.
13	Agent sends a captured AI prompt for inspection and applies the decision	The agent sends an AI prompt for inspection and then allows, masks or blocks it according to the company policy.
14	A registered company connects its first device end to end	Full journey check: a company registers, installs the agent on one machine and sees that device appear as Online.
15	First protected AI activity appears in the company audit log	After an employee sends an AI prompt, the company admin sees a record of it with sensitive values hidden.

Sprint outcome: A company signs up, verifies its email, logs in, sets one policy, installs the agent on one Linux machine, sees that device as Online, and sees its first inspected AI prompt.

02 Sprint 2 — Sensitive Data Protection & Full Linux Coverage

15 Aug - 28 Aug · Labels: DATA MASKING, ENDPOINT AGENT, SAAS

#	JIRA TITLE	DESCRIPTION
1	Detect common sensitive data (emails, names, cards, passwords, addresses)	Recognise everyday personal and company information inside AI prompts so it can be protected.
2	Detect IDs (CNIC, SSN, IBAN, phone, passport)	Recognise national and financial identity numbers used by customers in different countries.
3	Mask sensitive data in AI requests and responses, store masked records only	Replace sensitive values before the prompt reaches the AI provider, and keep only the masked version in company records.
4	Company admin chooses which sensitive data types to protect	The admin turns individual data categories on or off to match the company's own privacy rules.
5	Subscription plans page (Free / Pro / Enterprise)	Show the available plans with what each one includes so a customer can pick the right one.
6	Apply plan limits (maximum devices, monthly prompt quota)	Enforce the device count and monthly usage allowed by the company's plan.
7	Install the company security certificate so traffic can be inspected safely	Install and trust the security certificate on the employee machine so protected traffic can be checked without breaking the browser.
8	Linux: protect AI apps, developer tools and API traffic	Protect AI usage coming from installed applications, developer tools and scripts on Linux machines.
9	Linux: protect ChatGPT / Claude / Gemini websites	Protect what employees type into the ChatGPT, Claude and Gemini websites in their browser.
10	Send protected AI website activity to the company console	Website AI activity is reported to the company console so admins see it alongside application activity.
11	Linux: one-command install that starts protection automatically	IT installs the agent with a single command, and protection starts automatically and restarts with the machine.
12	Linux: reinstall without duplicate devices or wasted licences	Reinstalling on the same machine updates the existing device instead of creating a duplicate and consuming another licence.

#	JIRA TITLE	DESCRIPTION
13	Sensitive data typed into ChatGPT on Linux is masked before it leaves the machine	End-to-end check that real sensitive information typed into an AI website never reaches the provider.
14	A company that reaches its device limit cannot enrol more machines	End-to-end check that plan limits are enforced when IT tries to install beyond the purchased device count.

Sprint outcome: A company decides what sensitive data to protect, installs the agent on Linux machines with one command, and employee AI website and AI application activity is masked.

03 Sprint 3 — Company Console & Deployment Experience

29 Aug - 11 Sep · Labels: ADMIN DASHBOARD, AGENT UX, SAAS, SUPER ADMIN

#	JIRA TITLE	DESCRIPTION
1	Company security overview (protected prompts, masked data, blocked activity)	A dashboard showing how much AI activity was protected, what was masked, what was blocked and how this changes over time.
2	Company device fleet: Online / Offline / revoke access	List every company device with its status, and let the admin revoke a machine that should no longer be protected or trusted.
3	Company install credentials: view, copy and rotate	The admin views and copies the install token for new devices, and rotates it if it may have been shared or leaked.
4	Company settings (profile, preferences, protection configuration)	One place for the admin to manage company details, preferences and protection configuration.
5	Company plan and usage page (plan, quota used, quota remaining)	The admin sees the current plan, how many devices and prompts have been used, and how much allowance is left.
6	Super Admin customer overview (companies, plans, usage, status, fleet size)	The Super Admin reviews all customers in one view with their plan, usage, account status and number of protected devices.
7	Every install binds to its company so activity is attributed correctly	A device installed from a company's package reports only into that company, so activity is never attributed to the wrong customer.
8	Download agent page with a company-specific Linux package and steps	The admin downloads an installer already tied to their company, with clear installation instructions.
9	Guided deploy wizard: choose the operating system, install, confirm Online	A step-by-step wizard walks the admin through choosing an operating system, installing the agent and confirming the device is protected.
10	Device status for employees: Active / Disconnected / Blocked	Employees can see on their own machine whether protection is Active, Disconnected or Blocked.

Sprint outcome: Company admins and Super Admin run the service day to day — monitor security, manage devices, deploy agents and track plan usage — without developer assistance.


04 Sprint 4 — Windows Protection, Threat Defence & Paid Plans

12 Sep - 25 Sep · Labels: WINDOWS, THREAT SHIELDING, SAAS, BILLING

#	JIRA TITLE	DESCRIPTION
1	Password reset and account recovery	An admin who forgets their password can securely regain access without contacting support.
2	Detect dangerous AI prompts (jailbreak, prompt injection, policy violations)	Identify prompts that try to bypass AI safety controls or break company rules, and act on them.
3	Buy a paid plan: checkout	A customer upgrades from the plans page, pays, and their new plan limits apply automatically.
4	Super Admin billing view (paid plans, failed payments, plan status)	The Super Admin reviews paid subscriptions, plan status and failed payments across all customers.
5	Contact sales / Enterprise request form	A larger prospect submits their requirements so the sales team can follow up.
6	Windows: always-on protection that starts with the computer	The agent runs on Windows automatically from start-up and keeps running in the background without user action.
7	Windows: security certificate trust and AI application traffic protection	Trust the security certificate on Windows and protect AI usage from installed applications and developer tools.
8	Windows: protect ChatGPT / Claude / Gemini websites	Protect what employees type into AI websites from Windows browsers.
9	Windows: silent enterprise rollout via MSI, Group Policy or Intune	IT deploys the agent to many Windows machines at once, with no prompts for the employee.
10	Windows: signed installer and tray app showing protection status	Ship a trusted signed installer and a small status indicator so employees can confirm protection is active.
11	Single console view across Linux and Windows devices	The admin manages and monitors Linux and Windows machines together in one console.

Sprint outcome: The product serves Linux and Windows customers, blocks malicious AI prompts, and sells itself through self-service paid plans and enterprise enquiries.

Dependency Map Across Sprints



Feature Coverage Summary

AREA	SPRINT 1	SPRINT 2	SPRINT 3	SPRINT 4
Super Admin	List, suspend, activate companies	—	Customer overview and usage	Billing and subscription view
Company admin	Signup, login, first policy, first audit record	Choose protected data types, view plans	Dashboard, fleet, settings, usage, deploy wizard	Password recovery, cross-platform console
SaaS platform	Data isolation, company workspace	Plans and limits	Install credentials, usage tracking	Checkout, subscriptions, enterprise enquiries
Linux agent	Skeleton, enrolment, heartbeat, first inspection	Certificate trust, API and website protection, install, reinstall	Company-bound package, status indicator	Shared engine reused by Windows
Windows agent	—	—	—	Service, interception, websites, MSI, signed installer
Security engine	First inspection and audit entry	Sensitive data detection and masking	Security metrics	Threat detection